

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

Recurrent Neural Networks

Data Science Amsterdam

Rob Romijnders

Eindhoven, University of Technology

RomijndersRob@gmail.com

October 27, 2016

Overview

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

1 Why

2 How

3 What

4 Questions

GMail reply

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

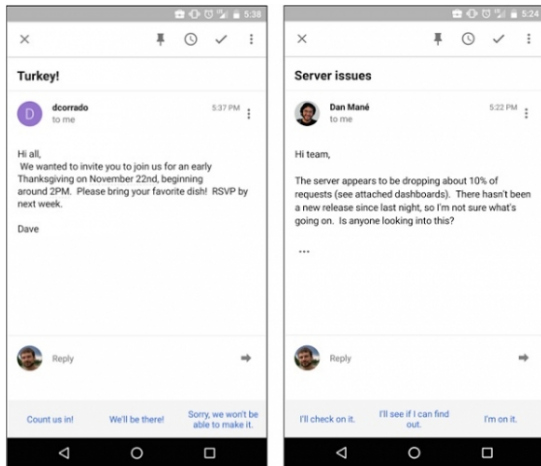


Figure: Photo: Greg Corrado, Google Research Blog

Apple siri

October 27,
2016

Rob
Romijnders

Why

How

What

Questions



Siri



Figure: Photo: cultofmac.com

Stock volatility prediction

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

Deep Learning Stock Volatility with Google Domestic Trends

Ruoxuan Xiong¹, Eric P. Nichols² and Yuan Shen^{*3}

¹Department of Management Science and Engineering, Stanford University

²Google Inc.

³Department of Physics, Stanford University

Generation

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

Data Science Amsterdam
Data Science Amsterdam
Data Science Amsterdam

Figure: Generating Sequences With Recurrent Neural Networks, Alex Graves
Music: github.com/hexahedria

Generation

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

THE MULTIVERSE —

Movie written by algorithm turns out to be hilarious and intense

For *Sunspring*'s exclusive debut on Ars, we talked to the filmmakers about collaborating with an AI.

ANNALEE NEWITZ · 6/9/2016, 12:30 PM



Sunspring, a short science fiction movie written entirely by AI, debuts exclusively on Ars today.

Figure: <http://arstechnica.com/the-multiverse/2016/06/an-ai-wrote-this-movie-and-its-strangely-moving/>

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

Deep learning

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions



¹Inception 5 (GoogLeNet)



Inception 7a

¹Going Deeper with Convolutions, [C. Szegedy et al, CVPR 2015]

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Deep learning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

RNN cell

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

$$h = \phi(W_{xh}x)$$

$$y = \phi(W_{hy}h)$$

$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1}) \quad y_t = \phi(W_{hy}h_t)$$

Python code

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

```
rnn = RNN()  
y = rnn.step(x)
```

```
class RNN:  
    def step(self, x):  
        self.h = np.tanh(np.dot(self.W_hh, self.h)  
                          + np.dot(self.W_xh, x))  
        y = np.dot(self.W_hy, self.h)  
        return y
```


Example

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Example

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

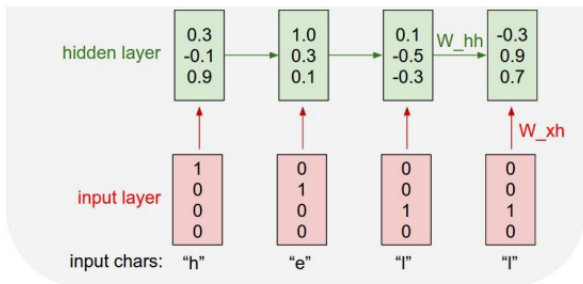


Figure: cs231n, Andrej Karpathy

Example

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

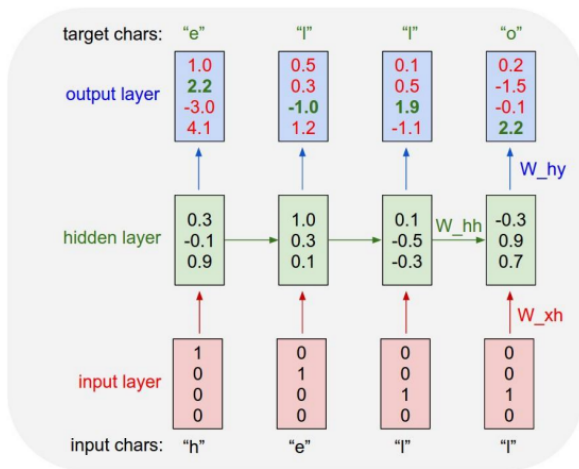


Figure: cs231n, Andrej Karpathy

Long Short-term memory

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i)$$

$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f)$$

$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o)$$

$$c_t = f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$$

$$h_t = o_t \tanh(c_t)$$

LSTM

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

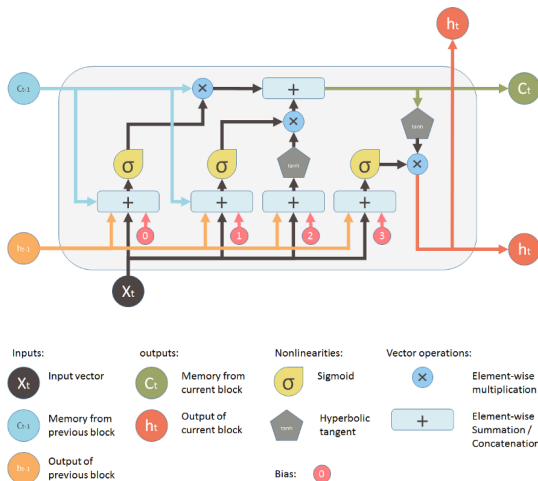


Figure: Diagram depicting LSTM block

Author: Shi Yan. Source: medium.com/@shiyan/

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Data

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: Classification of fractionated electrograms in epicardial mappings using a recurrent neural network, R Romijnders et al.

Data

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: Applying Deep Learning to Basketball Trajectories, R Shah, R Romijnders

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Data

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

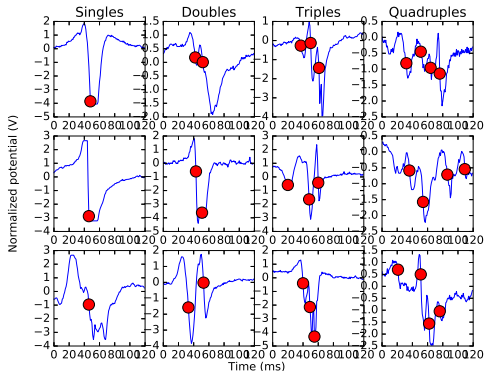


Figure: Classification of fractionated electrograms in epicardial mappings using a recurrent neural network, R Romijnders et al.

Annotation

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

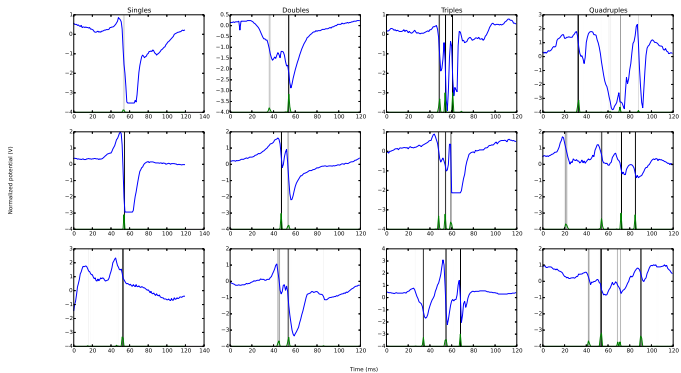


Figure: Classification of fractionated electrograms in epicardial mappings using a recurrent neural network, R Romijnders et al.

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

2D visualizations

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

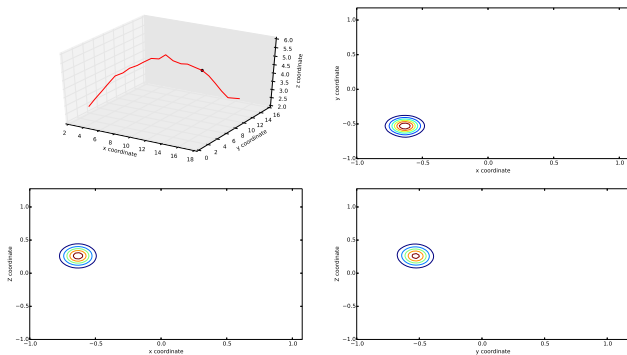


Figure: Applying Deep Learning to Basketball Trajectories, R Shah, R Romijnders

Sample

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: Applying Deep Learning to Basketball Trajectories, R Shah, R Romijnders

Biased sample

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: Applying Deep Learning to Basketball Trajectories, R Shah, R Romijnders

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

Captioning

October 27,
2016

Rob
Romijnders

Why

How

What

Questions



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."

Figure: Deep Visual-Semantic Alignments for Generating Image Descriptions, Andrej Karpathy

Architecture

October 27,
2016

Rob
Romijnders

Why

How

What

Questions




Figure: cs231n, Andrej Karpathy

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

Questions

Further reading

October 27,
2016

Rob
Romijnders

Why

How

What

Questions

- 1 The Unreasonable Effectiveness of Recurrent Neural Networks, Andrej Karpathy
- 2 Supervised sequence labelling with RNN, Alex Graves
- 3 Generating Sequences with RNN, Alex Graves
- 4 robromijnders.github.io